Mushroom Classification & Semi-supervised Learning

Group 2 Members:

Binaya Kumar Chaudhary Hemant Raj Singh Rohan Aryan-C0912902 Sani Asnain-C0906772 Shreya Baral (C0913115)



Table of Content



Business Case Evaluation



Data Identification



Data Acquisition and Filtering



Data Validation and Cleansing



Data Aggregation and Representation



Data Analysis



Model Building



Utilization of Analysis Results





"In 2022, sales of fresh and processed mushrooms amounted to approximately 694.47 million Canadian dollars, an increase from around 653.51 million the previous year."





Business Case Evaluation

Why classify mushrooms can be helpful?

- Risks of consuming Toxic Mushrooms
- The need for accurate and Efficient Testing
- Improved Food Safety for Consumers
- Reduced Healthcare costs from Mushroom Poisoning

Data Identification



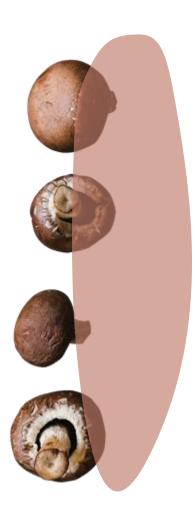
A mushroom containing 173 species (353 mushrooms per species).



Obtained from the UCI dataset repository.



Can be accessed at the following link:
https://archive.ics.uci.edu/dataset/848/secondary+mushroom+dataset.



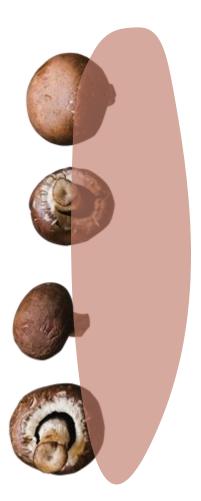
Dataset Features

Class information:

class: poisonous=p, edible=e (binary)

Variable Information: (n: nominal, m: metrical; nominal values as sets of values)

- 1. cap-diameter (m): float number in cm
- 2. cap-shape (n): bell=b, conical=c, convex=x, flat=f, sunken=s, spherical=p, others=o
- **3. cap-surface (n):** fibrous=i, grooves=g, scaly=y, smooth=s, shiny=h, leathery=l, silky=k, sticky=t, wrinkled=w, fleshy=e
- **4. cap-color(n):** brown=n, buff=b, gray=g, green=r, pink=p, purple=u, red=e, white=w, yellow=y, blue=l, orange=o, black=k
- **5. does-bruise-bleed (n):** bruises-or-bleeding=t, no=f
- **6. gill-attachment (n):** adnate=a, adnexed=x, decurrent=d, free=e, sinuate=s, pores=p, none=f, unknown=?
- 7. **gill-spacing (n):** close=c, distant=d, none=f
- 8. **gill-color(n):** see cap-color + none=f



Dataset Features

- 1. stem-height (m): float number in cm
- 2. stem-width (m): float number in mm
- **3. stem-root (n)**: bulbous=b, swollen=s, club=c, cup=u, equal=e, rhizomorphs=z, rooted=r
- 4. stem-surface (n): see cap-surface + none=f
- 5. **stem-color(n)**: see cap-color + none=f
- 6. veil-type (n): partial=p, universal=u
- 7. **veil-color (n)**: see cap-color + none=f
- **8.** has-ring(n): ring=t, none=f
- **9. ring-type (n)**: cobwebby=c, evanescent=e, flaring=r, grooved=g, large=l, pendant=p, sheathing=s, zone=z, scaly=y, movable=m, none=f, unknown=?
- 10. spore-print-color (n): see cap color
- **11.** habitat (n): grasses=g, leaves=l, meadows=m, paths=p, heaths=h, urban=u, waste=w, woods=d
- 12. season (n): spring=s, summer=u, autumn=a, winter=w

Data Acquisition and Filtering



Data source: CSV format



File size: below 5MB



Library used: pandas, sklearn, seaborn, numpy



Function used to load dataset: read_csv()

Splitting Data



```
[] #dividing the features in numerical and categorical types
   Target='class'
   features=[col for col in df.columns if col not in [Target]]
   numeric_features=['cap-diameter','stem-height','stem-width']
   categorical_features=[col for col in features if col not in numeric_features]
   print('the target variable is:',Target)
   print('='*90)
   print('numeric_features',numeric_features)
   print('lenght:',len(numeric_features))
   print('='*90)
   print('categorical_features:',categorical_features)
   print('lenght:',len(categorical_features))
   the target variable is: class
   numeric_features ['cap-diameter', 'stem-height', 'stem-width']
   lenaht: 3
   categorical_features: ['cap-shape', 'cap-surface', 'cap-color', 'does-bruise-or-bleed', 'gill-attachment', 'gill-spacing', 'gill-color', 'stem-root', 'stem-surface', 'stem-color', 'veil-type', 'veil-
   lenght: 17
```



Data Validation & Cleansing









	missing_percent
veil-type	94.797688
spore-print-color	89.595376
veil-color	87.861272
stem-root	84.393064
stem-surface	62.427746
gill-spacing	41.040462
cap-surface	23.121387
gill-attachment	16.184971
ring-type	4.046243
class	0.000000
stem-color	0.000000
habitat	0.000000
has-ring	0.000000
stem-width	0.000000
cap-diameter	0.000000
stem-height	0.000000
gill-color	0.000000
does-bruise-or-bleed	0.000000
cap-color	0.000000
cap-shape	0.000000
season	0.000000

We encountered large percentage of missing values in the data set, along with some special characters like '?, nan'

We took up imputation for handling the missing data

```
[ ] #replacing special characters with NA
    df.replace(['?', 'nan'],pd.NA,inplace=True)

[ ] #imputing nan values with mode
    for columns in df.columns:
        most_common_value=df[columns].mode()[0] #to use the first most occuring value, as there can be a tie between value occurance
    df[columns].fillna(most_common_value,inplace=True)

[ ] #checking if the fillna worked.
    df.isnull().sum()
```







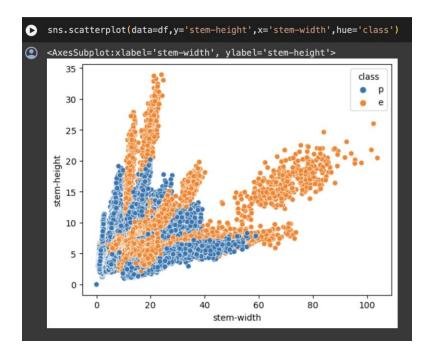
Data Normalization

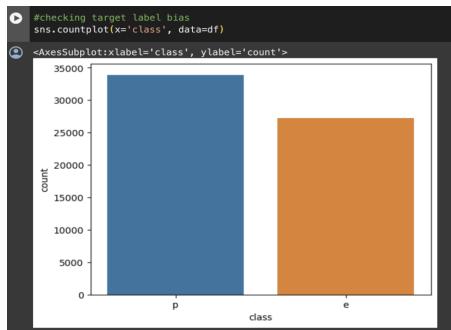
To pre-process the data for model building we took up the approach of using pipelines.

We created separate pipeline for categorical and numerical features, and integrated those pipelines using a column transformer

```
#converting the colums to lists:
numerical list=list(numeric features)
cateorical list=list(categorical features)
#pipeline for categorical features
pipeline categorical = Pipeline([
  ('onehot', OneHotEncoder(drop='first')),
#pipeline for numerical features
pipeline numerical = Pipeline([
  ('scaler', StandardScaler()),
#combined pipeline using column transformer
pipeline full = ColumnTransformer([
    ("numerical", pipeline numerical, numerical list),
    ("categorical", pipeline categorical, cateorical list),
```

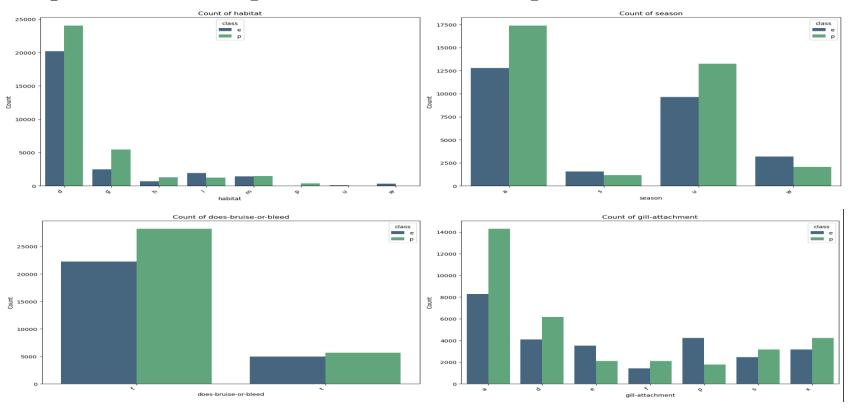


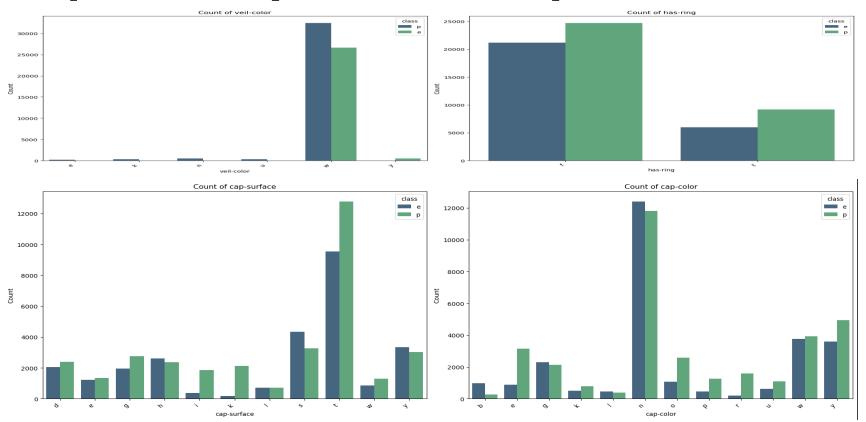






```
Count of cap-shape
                                                                                             14000 -
def plot_count(x,ax):
                                                                                                   р
        group = df.groupby([f'{x}','class'])['class'].count().reset_index(name='Count')
        sns.barplot(data=group, x=x, y='Count', hue='class', palette='viridis', ax=ax)
                                                                                             12000 -
        ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
        ax.set_title(f'Count of {x}')
                                                                                             10000
    cols = df.columns.tolist()
    fig, axes = plt.subplots(10, 2, figsize=(18, 6 * 11)
    for index,column in enumerate(cols[1:]):
        row = index // 2
        col = index % 2
        ax = axes[row,col]
        plot_count(column, ax)
                                                                                              2000
    plt.tight_layout()
    plt.show()
```





Semi-Supervised Learning

Semi-supervised approach, where characteristics of both supervised and unsupervised learning approaches are inhibited. This is a different approach to train the model, in this case a classifier. First with a smaller portion of labeled data and then a larger portion of unlabelled data. This approach pre-trains the model to understand the inherent data properly.

In our project, We used the last 20,000 records of our primary data set to fit into K-Means Clustering. For this we divided the data into labelled and unlabelled, trained the model on a small set and made predictions for a larger set, we then replaced these 20,000 records with 'Pseudo-labels' and used the now integrated data set for our other models.

Goal of using this approach is to uncover intricate relationship between the features with the target feature and pre-train the model for it, to have better results with the larger original dataset.



K-Means Clustering

- 20,000 records of primary data set to fit into K-Means Clustering
- We pre-processed the data using pipelines
- We then split the data into labelled and unlabelled segments

For semi-supervised approach, we made sure to have a smaller training set and have a larger testing set to have more predicted labels.

```
#using the last 20k records
df1=df.loc[41069:61069]
#shape of the segmented data
df1.shape
(20000, 21)
#x wil have labelled data which will be used to train the model
#y will have unlabelled data which will be used to test the model
#in semi supervised approach, training data is in smaller quantity compared to testing data
X=df1[features]
y=df1[Target]
print('the feature set shape is:',X.shape)
print('the Target set shape is:',y.shape)
the feature set shape is: (20000, 20)
the Target set shape is: (20000,)
```

K-Means Clustering



We did not require to use any additional technique to determine the number of clusters as we were dealing with a only 2 classes.

Shape of the data set after the split and preprocessing is shown in this step here.

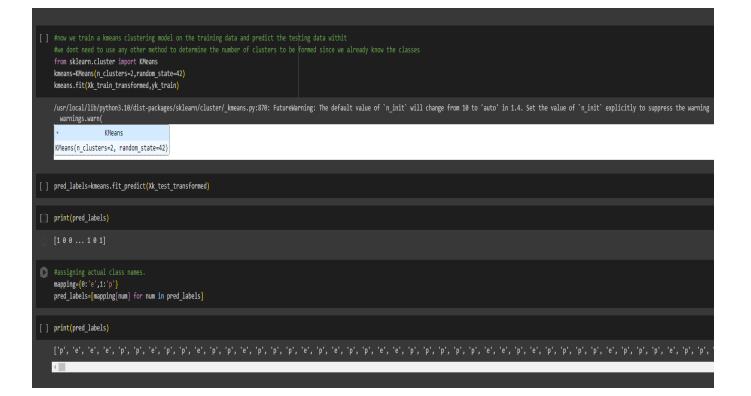
As we had high number of categorical variables. We ended up generating around 87 columns after pre-processing

Creating Labelled and Unlabelled Data Here: Xk_train and yk_train are labelled data and Xk_test is unlabelled Data #now we split the feature and the target set in 20:80 ratio using train test split #we dont need a y test to compare the pred vs actual in this case, as the whole purpose is to just train the model #and gather pseudo labels to be merged into the main dataset. from sklearn.model selection import train test split Xk train, Xk test, yk train, = train test split(X, y, test size=0.8, stratify=y) #stratified sampling based on the target print(f"Xk train.shape: {Xk train.shape}") print(f"Xk test.shape: {Xk test.shape}") print(f"yk_train.shape: {yk_train.shape}") Xk train.shape: (4000, 20) Xk test.shape: (16000, 20) yk_train.shape: (4000,) #pre-processing using pipelines. pipeline full.fit(Xk_train) Xk train transformed=pipeline full.fit transform(Xk train) Xk_test_transformed=pipeline_full.fit_transform(Xk_test) print(Xk train transformed.shape) print(Xk test transformed.shape) (4000, 87) (16000, 87)

K-Means Clustering



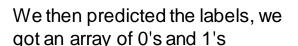






Data Aggregation and Representation

```
#merging the data with main dataset df1:
df_temp = pd.DataFrame({'class': pred_labels}, index=Xk_test.index)
df_20=pd.merge(Xk_test, df_temp, left_index=True, right_index=True, how='right')
#arranging the 'class' variable in the right order
df_20=df_20.set_index('class', append=True).reset_index(level=-1)
#merging Xk train and yk train:
merged_df=pd.concat([Xk_train, yk_train], axis=1)
#arranging class under the right column
merged df=merged df.set index('class', append=True).reset index(level=-1)
#merging df 20 amd merged df
final df=pd.concat([df 20,merged df])
#removing the last 20k records with actual lables and replacing it with pseudo lables.
df=df.iloc[:-20000]
#merging the pseudo lable dataset to main dataset
df=pd.concat([df,final df])
```



We mapped the data accordingly and integrated it into the main dataset





Model Building

KNearestNeighbours

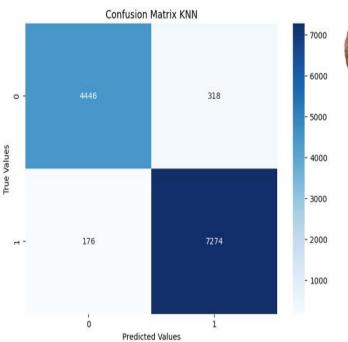
Random Forest Classifier

Decision Tree Classifier

Logistic Regression

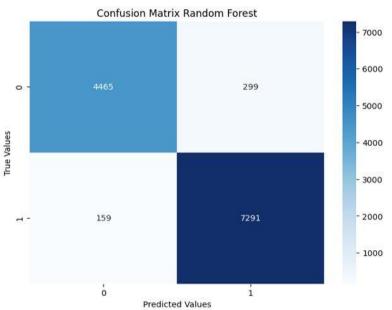


```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model selection import GridSearchCV
from sklearn.model selection import cross val predict
from sklearn.metrics import accuracy score, confusion matrix, roc auc score
knn classifier=KNeighborsClassifier()
knn classifier.fit(X train, y train)
y pred = knn classifier.predict(X test)
y probas trees = cross val predict(knn classifier, X train, y train, cv=4, method="predict proba")
y_tree_scores = y_probas_trees[:, 1] # score = proba of positive class
roc auc trees = roc auc score(y train,y tree scores)
test accuracy = accuracy score(y test, y pred)
print(f"Test Set Accuracy for model {knn_classifier} is {test_accuracy:.2f}")
Test Set Accuracy for model KNeighborsClassifier() is 0.96
```





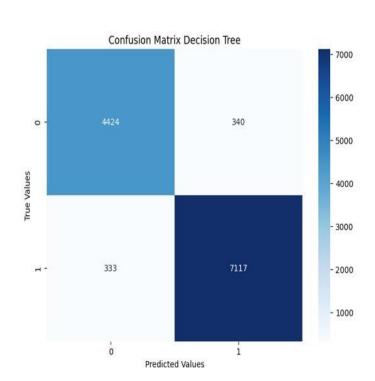
```
from sklearn.ensemble import RandomForestClassifier
rf_classifier = RandomForestClassifier(random state = 42)
rf classifier.fit(X train, y train)
y_pred = rf_classifier.predict(X_test)
y_probas_trees = cross_val_predict(rf_classifier,X_train,y_train, cv=4, method="predict_proba")
y_tree_scores = y_probas_trees[:, 1] # score = proba of positive class
roc_auc_trees = roc_auc_score(y_train,y_tree_scores)
test_accuracy = accuracy_score(y_test, y_pred)
print(f"Test Set Accuracy for model {rf_classifier} is {test_accuracy:.2f}")
Test Set Accuracy for model RandomForestClassifier(random_state=42) is 0.96
```





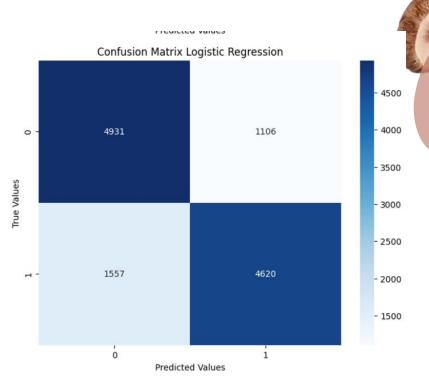


```
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import RandomizedSearchCV
clf = DecisionTreeClassifier(random_state=42)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
y_probas_trees = cross_val_predict(clf,X_train,y_train, cv=4, method="predict_proba")
y_tree_scores = y_probas_trees[:, 1] # score = proba of positive class
roc_auc_trees = roc_auc_score(y_train,y_tree_scores)
test_accuracy = accuracy_score(y test, y pred)
print(f"Test Set Accuracy for model {clf} is {test_accuracy:.2f}")
Test Set Accuracy for model DecisionTreeClassifier(random_state=42) is 0.94
```

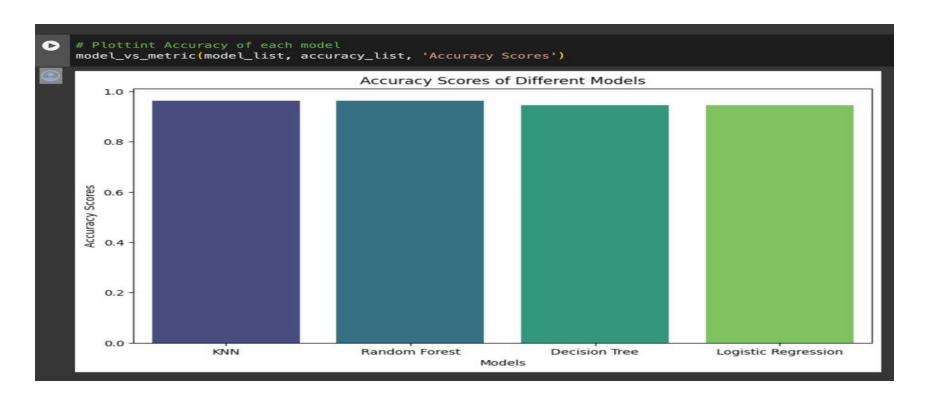




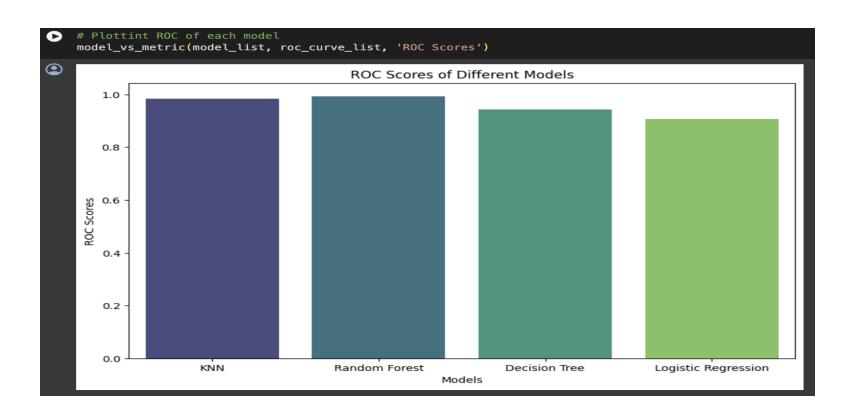
```
from sklearn.linear model import LogisticRegression
logreg model = LogisticRegression(max iter=10000)
logreg model.fit(X train, y train)
y pred = logreg model.predict(X test)
y_probas_trees = cross_val_predict(logreg_model,X_train,y_train, cv=4, method="predict_proba")
y_tree_scores = y_probas_trees[:, 1] # score = proba of positive class
roc_auc_trees = roc_auc_score(y_train,y_tree_scores)
accuracy = accuracy_score(y_test, y_pred)
print(f"Test Set Accuracy for model {logreg_model} is {test_accuracy:.2f}")
Test Set Accuracy for model LogisticRegression(max iter=10000) is 0.94
```



Accuracy Scores of Different Models



Roc Scores of Different Models



Utilization of Analysis

Data-Driven Decision-Making:

instead of randomly guessing a mushroom to be edible or not, you can take a data backed decision. The model generates valuable insights into the factors influencing the classification of mushrooms. These insights can be leveraged for data-driven decision-making, such as optimizing cultivation practices or adjusting product portfolios.

Market Expansion and Consumer Education:

A reliable classification model can facilitate market expansion by assuring consumers of the safety of the products. It also provides an opportunity for businesses to educate consumers about the importance of choosing safe mushrooms.

Market Differentiation:

A business that can guarantee the safety of its mushroom products through an advanced classification system may gain a competitive edge in the market. It can be positioned as a brand that prioritizes consumer well-being.

Crisis Management:

In the event of a contamination or poisoning incident, having a classification model can enable businesses to quickly identify affected batches, initiate recalls, and communicate with consumers, mitigating the impact on both public safety and the brand.

Thanks!

Any Questions?

